FreqODES: Frequency-Domain Guided Online Domain Adaptation with Expert-in-the-Loop Active Learning

Replication and extension of ODES (MICCAI 2025)

Group A12

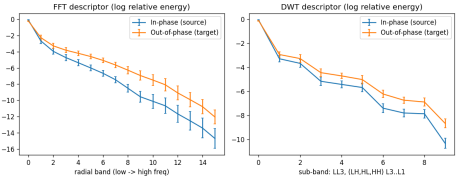
Vishnu Vardhan (CB.SC.U4AIE23043) Nithya Sai (CB.SC.U4AIE23021)

22AIE437 Medical Image Processing
B.Tech CSE (AI), Amrita School of Artificial Intelligence

November 7, 2026

Problem and motivation

- Abdominal organ segmentation (liver, kidneys, spleen) from MRI supports surgical planning, volumetry and follow-up.
- A network trained on one scanner / sequence (*source*) degrades on another (*target*): **domain shift**.
- Example: CHAOS T1-DUAL **in-phase** $\rightarrow$ **out-of-phase**: same patients, same organs, yet Dice of our source model drops from training level to **54.97%**.
- Retraining per hospital needs costly expert labels; unsupervised test-time adaptation relies on noisy pseudo-labels.
- **Online active domain adaptation**: adapt patient-by-patient, asking an expert to label only a tiny budget of pixels.



Same anatomy, different MR sequence: OOP has more high-frequency energy (FFT bands, left; DWT sub-bands, right).

Dataset, task and objectives

Dataset: CHAOS challenge [2] (MR, T1-DUAL)

- 20 patients, 647 axial slices (26–50 per patient), 256×256 / 288×288, resized to 256^2
- Labels: liver, right kidney, left kidney, spleen (+ background) = 5 classes
- **Source**: in-phase (IP), all 20 volumes, used for training
- **Target**: out-of-phase (OOP), the 20 volumes arrive as an *online stream*, one patient = one batch (paper protocol, no train/test split on target)

Task: 2-D multi-class **segmentation**

Objectives

1. Reproduce ODES [1] on CHAOS IP→OOP.
2. Replace its BN-statistics image-pruning score with an interpretable **FFT / DWT frequency-domain** shift score (FreqODES).
3. Compare at equal annotation budget ($K = 10, 50, 100\%$ slices, $b = 1\%$ pixels).

Metrics: Dice (DSC), Jaccard (IoU) per organ and mean; adaptation time; paired Wilcoxon test over the 20 patients.

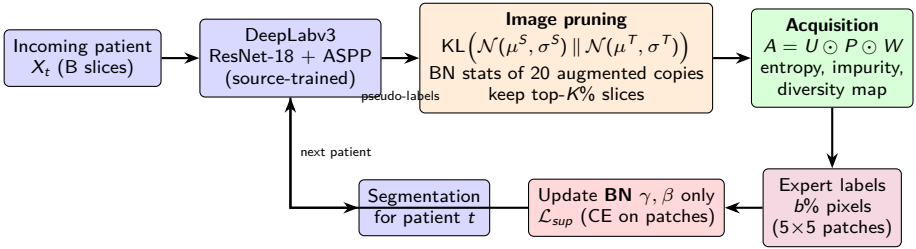
Base paper

ODES: Online Domain Adaptation with Expert Guidance for Medical Image Segmentation

M. S. Islam, S. Nag, A. Dutta, S. M. Ahmed, F. F. Niloy, S. Bera, A. K. Roy-Chowdhury
University of California, Riverside — **MICCAI 2025** [1]. Code: `github.com/ShazidAraf/ODES`

- First active-learning based *online* UDA framework for medical segmentation: source-free, storage-free, one pass per batch.
- Four stages per incoming patient: **infer** → **prune images** → **acquire patch labels** → **update BN layers**.
- Contributions: (i) **image pruning** by BN-statistics KL divergence; (ii) **diversity-aware acquisition** (entropy × regional impurity × farthest-first diversity map).
- Evaluated on CHAOS IP→OOP, CHAOS→DUKE, BMC→RUNMC; beats TENT, CoTTA, SaTTCA, ...

ODES architecture



- **Backbone**: DeepLabv3 [3], ResNet-18 encoder (ImageNet init), ASPP, low-level skip, 5-class head.
- **Pruning**: each slice is augmented (flip / rotation), passed through a copy of the net; BN statistics of the last 9 BN layers are compared with source running statistics (Eq. 1 of [1]).
- **Acquisition**: U = patch-mean entropy, P = regional impurity [4], W = Gaussian map around farthest-first-traversal picks (spatial $\times$ feature).

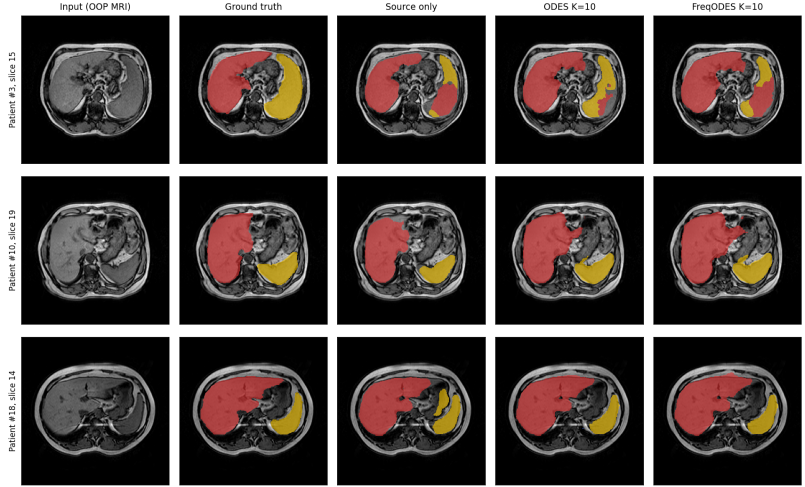
Replicated vs. reported results (CHAOS IP $\rightarrow$ OOP, $b = 1\%$)

Table 1: Dice (%) per organ; Jaccard (IoU, %) mean. Paper values from Table 1 of [1].

	Setting	Liver	L. Kidney	R. Kidney	Spleen	**Mean DSC**	**Mean IoU**
Source only	Paper	87.77	37.97	18.92	67.31	52.99	–
	Ours	85.89	29.84	33.74	70.43	54.97	42.74
ODES K=100	Paper	88.90	72.41	74.55	78.56	78.61	–
	Ours	87.89	76.90	75.28	78.00	**79.52**	66.96
ODES K=50	Paper	89.13	72.30	74.43	78.36	78.56	–
	Ours	88.02	76.66	75.71	77.89	**79.57**	67.00
ODES K=10	Paper	88.41	71.45	70.22	77.12	76.80	–
	Ours	87.88	75.19	74.18	78.06	**78.83**	66.24

- Source-only and adapted results are within $\approx$2 DSC of the paper; the adaptation gain (54.97$\rightarrow$79.52, $\approx +25$) and the K trend ($100 \approx 50 > 10$) are reproduced.
- No source checkpoint or training recipe is released, so we trained our own source model; organ-level differences (e.g. kidneys) come from that.
- Paper reports no Jaccard; ours is measured per volume.

Qualitative results: input, ground truth, predicted masks



Red: liver, green: right kidney, blue: left kidney, yellow: spleen. Three target (OOP) patients, slice with largest organ area.

Code structure and training settings

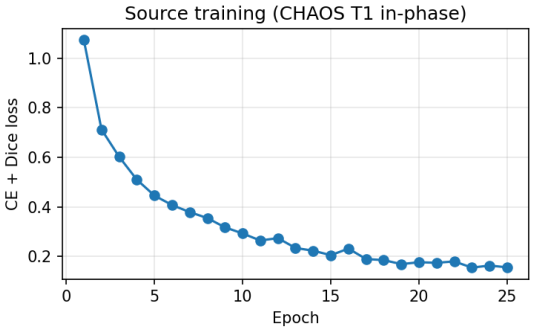
```
prepare_chaos.py    DICOM/PNG -> (D,H,W) tif stacks
dataloader/         Stack_Loader: one patient = batch
model/              DeepLabv3 (ResNet-18 + ASPP)
train_source.py     source training (CE + Dice)
methods/ODES.py     pruning, acquisition, BN update
methods/freq_shift.py  NEW: FFT / DWT shift score
FFT_weight_2.py     diversity map (missing upstream,
                    re-implemented from the paper)
demo_ODES.py        online loop + Dice / IoU eval
run_all.sh          21-run experiment sweep
make_report.py      tables, plots, LaTeX numbers
```

Upstream repo had no training script, no checkpoint and a missing module; we added these and fixed CPU/GPU and a stale-variable bug.

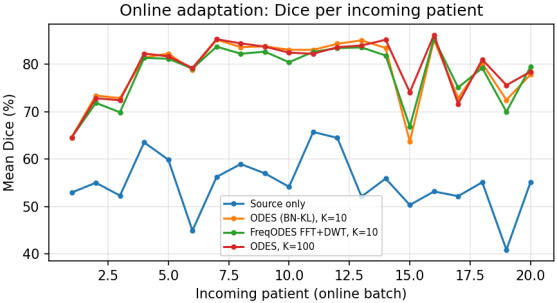
Table 2: Training / adaptation settings

Framework	PyTorch 2, timm
Source optimiser	Adam, lr 10^{-3}, wd 10^{-4}, cosine
Batch size / epochs	8 / 25
Loss	CE + Dice
Augmentation	flips, 90° rot., intensity jitter
Adaptation optimiser	Adam, lr 10^{-2}, 1 step / patient
Updated params	BN γ, β only
Budget	$b = 1\%$ px, $K \in \{10, 50, 100\}\%$
Hardware	NVIDIA Tesla T4 (Colab)
Seed	1337 (all runs)

Training log and curves



Source training loss (25 epochs, 647 IP slices): $1.07 \rightarrow 0.16$.



Online adaptation: mean Dice of each incoming target patient.

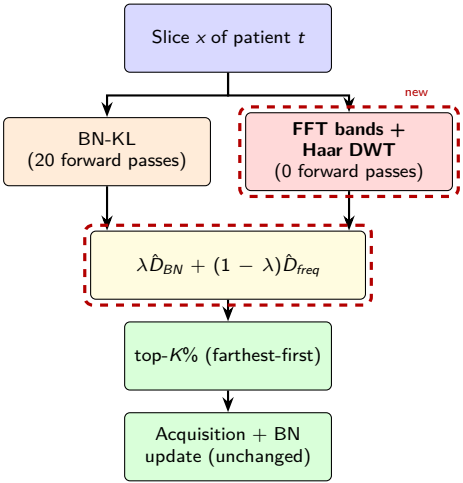
Expert-guided BN updates lift every patient above source-only.

What we changed: frequency-domain image pruning

Hypothesis: MR sequence / scanner shift is directly visible as a change in spatial-frequency energy, so it can be detected from the image itself, without the network.

- **FFT descriptor**: log fraction of energy in 16 concentric radial bands of $|\mathcal{F}(x)|^2$.
- **DWT descriptor**: log relative energy of LL_3 and LH/HL/HH sub-bands of a 3-level Haar DWT.
- Reference μ_f, σ_f from 200 random source slices.
- Shift score $D_{freq}(x) = \frac{1}{d}\sum_f \left(\frac{\phi_f(x)-\mu_f}{\sigma_f}\right)^2$
- **Fused**: $D = \lambda \hat{D}_{BN} + (1-\lambda)\hat{D}_{freq}$

Only Stage 3 changes; backbone, acquisition, update and budgets are identical $\Rightarrow$ fair comparison.



Does the frequency score detect the shift?

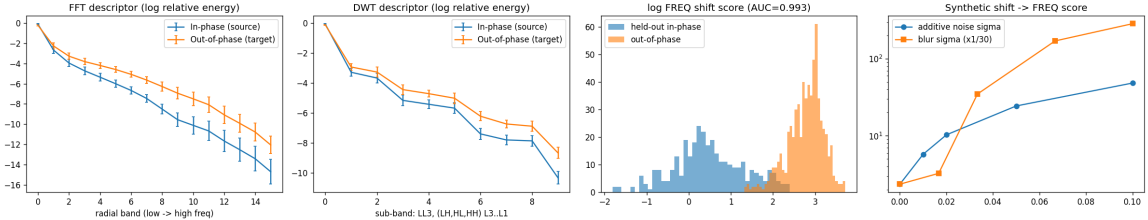


Table 3: Separating held-out IP slices (10 patients) from OOP slices with the shift score alone.

Descriptor	FFT	DWT	FFT+DWT
ROC-AUC	**0.994**	0.969	0.993

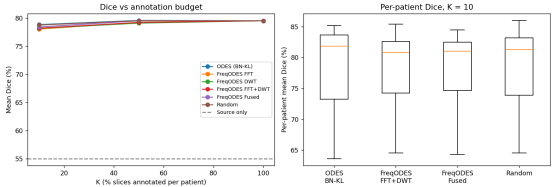
OOP slices carry systematically more high-frequency energy; the score also grows monotonically with synthetic noise and blur.

FreqODES v1 (frequency score alone) vs. ODES, single seed

Table 4: Same data, model, seed and budget ($b = 1\%$). DSC / IoU in %, time = full online pass over 20 patients on a T4. p: paired Wilcoxon vs. ODES at the same K.

	$K = 10\%$ of slices				$K = 50\%$ of slices			
Pruning score	DSC	IoU	p	Time (s)	DSC	IoU	p	Time (s)
BN-KL (ODES, replicated)	**78.83**	66.24	–	167	79.57	67.00	–	171
FFT (FreqODES)	78.02	65.03	0.005	107	79.24	66.61	0.064	115
DWT (FreqODES)	78.15	65.28	0.027	109	79.09	66.43	0.030	114
FFT+DWT (FreqODES)	78.20	65.27	0.091	109	79.20	66.52	0.044	114
Fused, λ=0.5	78.40	65.45	0.159	169	**79.50**	66.92	0.687	173
Random	78.75	66.04	0.334	107	79.52	67.00	1.000	113

All slices ($K = 100$): DSC 79.52, IoU 66.96 Source only: DSC 54.97, IoU 42.74



Ablations: pruning $\times$ diversity weighting, and fusion weight

Table 5: Mean DSC at $K = 10$ (cf. Table 3 of [1]).

Image pruning	D.W. off	D.W. on
BN-KL (ODES)	78.73	78.83
FFT+DWT (ours)	78.32	78.20
Random	78.96	78.75
Paper (BN / random)	75.88 / 74.32	76.80 / 75.41

Table 6: Fusion weight λ ($K = 10$).

λ	Mean DSC
0 (pure frequency)	78.20
0.25	79.09
0.5	78.40
0.75	78.93
1 (pure BN = ODES)	78.83

Mixing a little BN information into the frequency score ($\lambda = 0.25$) gives the best $K = 10$ result (79.09 vs. 78.83, not significant, $p = 0.75$); diversity weighting changes DSC by ≤ 0.2 in our runs, much less than reported in the paper.

Analysis: why v1 alone did not beat ODES (and what it led to)

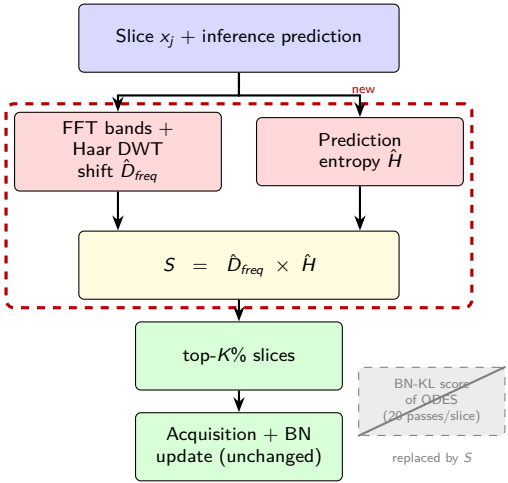
- **Where it helps**: best shift *detector* (AUC 0.994), interpretable, needs no forward passes $\Rightarrow$ full online pass 109 s vs. 167 s ($\approx$ 35% faster). The fused score matches ODES at both budgets (79.50 vs. 79.57 at $K = 50$, $p = 0.687$) and with $\lambda = 0.25$ reaches 79.09 at $K = 10$.
- **Where it hurts**: pure frequency scores trail ODES by 0.6–0.8 DSC at $K = 10$ (FFT: -0.81, $p = 0.005$) and by 0.3–0.5 at $K = 50$.
- **Mismatch between objective and score**: the most frequency-shifted slices are the volume extremes (few or tiny organs, mostly background); labelling them teaches the BN layers little. BN-KL measures what *the network* finds unusual, which correlates better with segmentation error.
- **Small, homogeneous shift**: IP$\rightarrow$OOP shift is global, so nearly every OOP slice is "shifted"; ranking among them is what matters, and random pruning is already strong (78.75, $p = 0.334$).
- **Fix**: a slice is worth labelling if it is *shifted* **and** the model is *unsure* about it $\Rightarrow$ multiply the frequency score by model uncertainty (next slide), and evaluate over 3 seeds.

FreqODES-U: frequency shift × model uncertainty

Slice score

$$S(x_j) = \underbrace{\hat{D}_{freq}(x_j)}_{\text{FFT+DWT shift}} \times \underbrace{\hat{H}(x_j)}_{\text{mean pixel entropy}}$$

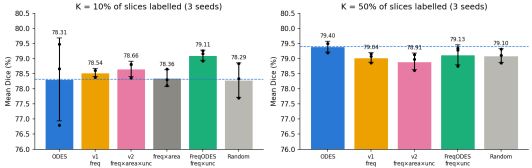
- $\hat{D}_{freq}$: v1 frequency score; $\hat{H}$: mean entropy of the softmax prediction of the slice; each normalised by its batch maximum.
- $\hat{H}$ reuses the inference prediction ODES already makes ⇒ still **0 extra forward passes**.
- Pre-registered design (v2): $\hat{D}_{freq} \times$ organ area $\times \hat{H}$, with two ablations (drop one factor each). The ablation $\hat{D}_{freq} \times \hat{H}$ was best; we report *all* variants.
- Protocol: 3 seeds (1, 2, 12345) × 20 patients; paired Wilcoxon over the 60 patient-runs.



FreqODES-U results: 3 seeds, same model and budget

Table 7: Mean Dice (%) ± std over 3 seeds; Δ and p: paired Wilcoxon vs. ODES over 60 patient-runs; time: full online pass (T4).

	$K = 10\%$ of slices				$K = 50\%$ of slices		
Selection score	DSC	Δ / p	IoU	Time	DSC	Δ / p	IoU
ODES (BN-KL)	78.31±1.38	–	65.51	154 s	**79.40**±0.18	–	66.90
v1: frequency	78.54±0.13	+0.22 / 0.908	65.63	104 s	79.04±0.15	-0.36 / 0.000	66.36
v2: freq × area × unc.	78.66±0.26	+0.35 / 0.199	65.84	105 s	78.91±0.27	-0.49 / 0.000	66.18
Ablation: freq × area	78.36±0.28	+0.04 / 0.700	65.57	104 s	–	–	–
FreqODES-U: freq × unc.	**79.11**±0.17	**+0.80 / 0.004**	**66.43**	104 s	79.13±0.33	-0.27 / 0.018	66.44
Random	78.29±0.57	-0.02 / 0.232	65.40	102 s	79.10±0.23	-0.30 / 0.001	66.43



At $K = 10\%$ FreqODES-U beats ODES on every seed (78.93, 79.16, 79.26 vs. 76.80, 78.67, 79.48), 8× more stable, ≈33% faster. At $K = 50\%$ ODES stays best.

Conclusion

- **Replication**: ODES reproduced on CHAOS IP→OOP after completing the public code: 79.52 DSC with all slices annotated (paper 78.61), adaptation gain ≈25 DSC as in the paper.
- **v1 – frequency shift score**: an interpretable FFT/DWT detector of domain shift (AUC 0.994) that needs no forward passes (≈35% cheaper), but on its own it does not beat ODES: a good *shift* detector is not automatically a good *annotation-value* detector.
- **FreqODES-U – frequency × uncertainty**: **79.11 vs. 78.31 DSC at** $K = 10\%$ (Δ+0.80, $p = 0.004$, 3 seeds), about 8× more stable and ≈33% faster; at $K = 50\%$ ODES stays best (79.40 vs. 79.13).
- **Caveats**: best variant of a pre-planned ablation; one dataset pair; 20 patients; 3 seeds.
- **Future work**: CHAOS→DUKE and BMC→RUNMC, more seeds and stream orders, organ-ROI spectra.

Team contributions

Both members took part in every stage and can explain the whole project; the main responsibilities were:

Vishnu Vardhan (CB.SC.U4AIE23043)

- Environment, CHAOS download and preprocessing
- Completing the ODES code: training script, missing `FFT_weight_2` module, bug fixes
- Source-model training and ODES replication
- Implementing the FFT/DWT score and FreqODES-U
- GPU experiments on Colab, GitHub repository

Nithya Sai (CB.SC.U4AIE23021)

- Literature survey and project proposal
- Frequency analysis of IP vs. OOP (spectra, AUC)
- v2 design, ablation plan and 3-seed evaluation
- Statistical analysis (Wilcoxon), figures and results page
- Slides, Elsevier report and viva preparation

References

[1] M. S. Islam, S. Nag, A. Dutta, S. M. Ahmed, F. F. Niloy, S. Bera, A. K. Roy-Chowdhury. ODES: Online domain adaptation with expert guidance for medical image segmentation. *MICCAI*, 2025.

[2] A. E. Kavur et al. CHAOS challenge – combined (CT-MR) healthy abdominal organ segmentation. *Medical Image Analysis* 69:101950, 2021.

[3] L.-C. Chen, G. Papandreou, I. Kokkinos, K. Murphy, A. L. Yuille. DeepLab: semantic image segmentation with deep convolutional nets, atrous convolution, and fully connected CRFs. *IEEE TPAMI* 40(4):834–848, 2018.

[4] B. Xie, L. Yuan, S. Li, C. H. Liu, X. Cheng. Towards fewer annotations: active learning via region impurity and prediction uncertainty for domain adaptive semantic segmentation. *CVPR*, 2022.

[5] D. Wang, E. Shelhamer, S. Liu, B. Olshausen, T. Darrell. Tent: fully test-time adaptation by entropy minimization. *ICLR*, 2021.

[6] Y. Yang, S. Soatto. FDA: Fourier domain adaptation for semantic segmentation. *CVPR*, 2020.

[7] S. G. Mallat. A theory for multiresolution signal decomposition: the wavelet representation. *IEEE TPAMI* 11(7):674–693, 1989.

Thank You

Questions?

Group A12 ▪ Vishnu Vardhan (CB.SC.U4AIE23043) ▪ Nithya Sai (CB.SC.U4AIE23021)

https://github.com/Vishnu868/FreqODES